

head/tail counters in each row

satellite

read/told? wraparound?

18,744,551,610 = 1,845 × 10,943,709,073 = 1,845 × 10¹²

1.82x10⁸) items/s (second through fourth quarter 1987, 1.82x10⁸ items/s) was 1.82x10⁸ items/s (first quarter 1988).

$$\bullet \quad 1.845 \times 10^{19} / 1.82 \times 10^8$$

$$= 100,809,041,482 \text{ s}$$

$$= 3,197 \text{ years}$$

Optimisations

Monotonic head/tail counters

- Will head/tail wrap around?
 - $2^{64} = 18,446,744,073,709,551,616 = 1.845 \times 10^{19}$
 - queue throughput was $182,987,000 (1.82 \times 10^8)$ items/second
 - $1.845 \times 10^{19} / 1.82 \times 10^8$
 $= 100,809,041,482\text{s}$
 $= 3,197$ years

```
template<typename T>
class drow_queue_v2
{
public:
    drow_queue_v2 (size_t capacity_)
        : capacity (std::bit_ceil (capacity_))
    {}

    bool try_push (const T&);

    bool try_pop (T&);

private:
    size_t capacity = 0;
    std::vector<T> data { std::vector<T> (capacity) };
    std::atomic<size_t> head { 0 }, tail { 0 };

};
```

```
bool try_push (const T& v)
{
    size_t current_tail = tail.load();
    size_t current_head = head.load();

    size_t size = current_tail - current_head;

    if (size >= (capacity - 1)) // full
        return false;

    size_t index = current_tail & (capacity - 1);
    data[index] = v;
    tail.store (current_tail + 1);

    return true;
}
```

```
bool try_pop (T& v)
{
    size_t current_head = head.load();
    size_t current_tail = tail.load();

    if (current_head == current_tail) // empty
        return false;

    size_t index = current_head & (capacity - 1);
    v = data[index];
    head.store (current_head + 1);

    return true;
}
```